

Edge-AI Malicious Domain and URL Detection for IoT Gateway Security: A Lightweight Random Forest Approach

Phan Luu Tung*, Nguyen Nhat Huy*, Nhu Gia Nguyen *
 *Department of Computer Science
 International School
 Duy Tan University
 Da Nang 550000, Vietnam
 Email: phanluutung@dtu.edu.vn

Abstract—The rapid growth of Internet of Things (IoT) deployments has significantly increased the exposure of edge gateways to malicious domains and URLs, enabling phishing, malware delivery, and command-and-control communications. Existing cloud-based detection solutions suffer from high latency, privacy concerns, and limited applicability in resource-constrained environments. This paper presents an edge-based malicious URL detection system designed specifically for IoT gateway environments.

We evaluated five machine learning algorithms (Random Forest, Logistic Regression, SVM, XGBoost, Neural Networks) using a 31-dimensional feature set combining URL lexical analysis, DNS behavior patterns, SSL certificate metadata, and domain registration information. Random Forest achieved the best balance of performance and efficiency with 72.30% accuracy and 0.7089 F1-score while meeting strict edge deployment constraints.

Experimental evaluation using 100,000 URLs revealed significant performance variations across algorithms. Accuracy ranged from 63.6% to 72.4% with F1-scores between 0.494 and 0.709, while **model-only throughput** varied from **112 to 408 samples per second**. Feature importance analysis showed URL structure characteristics contribute 51.5% to classification decisions and domain metadata accounts for 33.8%. The Random Forest implementation requires 3.5 MB memory and achieves **10.50 ms end-to-end offline latency** (including feature extraction and model prediction), demonstrating practical feasibility for IoT gateway deployment.

Index Terms—Artificial Intelligence, Edge AI, machine learning comparison, feature importance analysis, IoT security, malicious URL detection, Random Forest, XGBoost, cybersecurity

I. INTRODUCTION

The proliferation of Internet of Things (IoT) devices has transformed modern cyber-physical systems, enabling smart homes, industrial automation, and intelligent transportation. However, the rapid expansion of IoT ecosystems has also introduced significant security risks, particularly through malicious domains and URLs used for phishing attacks, malware distribution, and command-and-control communications. IoT gateways, which aggregate and forward traffic from numerous devices, have become attractive attack targets due to their central role in network communication [1].

Traditional malicious URL detection approaches predominantly rely on cloud-based analysis, where traffic is forwarded

to centralized servers for inspection. While effective in high-resource environments, such solutions suffer from several limitations when applied to IoT systems. First, the additional network latency introduced by cloud communication can delay threat mitigation, making real-time protection infeasible. Second, transmitting raw traffic data raises privacy concerns, especially in sensitive environments such as smart healthcare or industrial networks. Finally, continuous cloud connectivity cannot be assumed in many edge deployments [2].

Recent advances in Artificial Intelligence (AI) and machine learning have enabled the development of sophisticated cybersecurity solutions capable of detecting complex attack patterns. However, deploying AI models at the network edge requires careful consideration of computational constraints, memory limitations, and energy efficiency. The emergence of Edge-AI paradigms presents an opportunity to combine the intelligence of AI-driven detection with the low-latency requirements of edge computing [3]. This convergence is particularly relevant for IoT security, where real-time threat detection must be balanced against resource limitations [4].

Recent advances in edge computing and lightweight machine learning have enabled security mechanisms to be deployed closer to data sources [3]. Edge-AI approaches allow real-time analysis while reducing latency and preserving data locality. However, designing accurate and efficient detection systems for resource-constrained edge devices remains challenging. Deep learning models, while powerful, often incur high computational and memory overhead, limiting their practicality for edge deployment [5].

In this paper, we propose a comprehensive Edge-AI-based malicious domain and URL detection framework with systematic artificial intelligence methodology evaluation specifically designed for IoT gateway environments. The system employs rigorous comparison of five AI approaches to establish **empirical basis**: Random Forest [11], Logistic Regression, Support Vector Machine (SVM), Extreme Gradient Boosting (XGBoost), and Neural Networks. Our hybrid feature extraction strategy integrates URL lexical features [6], DNS query behavior [7], SSL certificate metadata [8], and domain registration information into a 31-dimensional framework that enables **practically viable** AI classification performance suitable for cybersecurity applications.

Although Random Forest is a classical ML model, this work demonstrates how careful feature engineering and system-level optimization enable AI models to operate efficiently in edge-constrained IoT environments. The comprehensive AI evaluation reveals significant performance trade-offs (**112-408 samples/sec model-only throughput variation**) critical for informed edge deployment decisions. The system is fully containerized and exposed via a RESTful API to support real-time AI-driven threat detection.

Research Gap. While malicious URL and domain detection has been extensively studied using lexical, DNS, and learning-based approaches, practical deployment on IoT gateways remains insufficiently addressed. Specifically, existing works suffer from three critical limitations: (i) **inadequate deployment-level evaluation** - many studies focus on offline predictive performance without reporting deployment-relevant metrics such as memory footprint, feature extraction overhead, and end-to-end latency/throughput under realistic edge hardware constraints, (ii) **inconsistent multi-model comparison** - different studies evaluate single models or model families using different datasets, feature sets, and measurement protocols, making it difficult to derive actionable model-selection guidelines for resource-constrained environments, and (iii) **insufficient feature necessity analysis** - while heterogeneous feature sets combining URL lexical, DNS behavioral, SSL metadata, and WHOIS information are commonly proposed, there is limited quantitative evidence on which feature groups are essential versus optional when computational resources and data availability are constrained at the gateway level. Therefore, there is a clear need for a systematic edge-AI pipeline with controlled multi-model comparison, quantified deployment trade-offs (accuracy vs. latency/memory/throughput), and evidence-based feature importance analysis to enable informed design decisions for practical IoT gateway security systems.

The main contributions of this work are summarized as follows:

- We establish a systematic AI algorithm evaluation framework for edge computing environments, comparing five distinct approaches (Random Forest, Logistic Regression, SVM, XGBoost, Neural Networks) with comprehensive performance analysis across accuracy, latency, memory, and throughput dimensions.
- We develop a 31-dimensional feature engineering methodology that enables AI approaches to achieve **practically viable** cybersecurity performance for edge deployment, with Random Forest achieving 72.30% accuracy and 0.7089 F1-score, with quantified feature importance analysis revealing URL lexical features (51.5%) and domain metadata (33.8%) as primary contributors.
- We demonstrate **empirically-grounded** AI model selection for edge deployment, showing Random Forest provides a **reasonable balance** of cybersecurity performance (72.30% accuracy, 0.7089 F1-score), **acceptable** end-to-end latency (10.50ms), and deployment flexibility for resource-constrained IoT environments.
- We conduct rigorous AI methodology evaluation including feature importance analysis, ablation studies, and

comparative performance benchmarking, establishing an **experimental basis** for edge-optimized AI security systems with quantified deployment trade-offs.

II. RELATED WORK

Malicious domain and URL detection has been extensively studied using a variety of approaches, including signature-based methods, machine learning techniques, and deep learning models. Early works relied on blacklist-based detection and handcrafted rules, which are effective for known threats but fail to generalize to previously unseen or rapidly evolving attacks [9].

Machine learning-based approaches have demonstrated improved generalization by leveraging lexical features extracted from URLs, such as length, entropy, character distribution, and token patterns [6], [10]. Random Forest, Support Vector Machines, and Gradient Boosting models have been widely adopted due to their robustness and interpretability [11], [12]. However, many existing studies focus primarily on offline analysis and do not consider deployment constraints in real-world systems [13].

DNS-based detection methods analyze query behavior, such as query frequency, NXDOMAIN ratios, and time-to-live (TTL) characteristics, to identify anomalous domain activities [7], [14]. These techniques are particularly effective against fast-flux and algorithmically generated domains [15], [16]. Nevertheless, DNS-only approaches may suffer from limited discriminative power when attackers mimic benign query patterns.

Recent advances in deep learning have introduced convolutional neural networks, recurrent neural networks, and Transformer-based models for malicious URL detection [17], [18]. While these models achieve high accuracy, their computational and memory requirements often exceed the capabilities of resource-constrained edge devices. Similarly, graph-based approaches using Graph Neural Networks (GNNs) [19] model relationships between domains, IP addresses, and DNS resolvers, providing strong contextual understanding at the cost of increased complexity.

Edge-AI security solutions have gained attention as a means to reduce latency and preserve data locality [3], [20]. However, many existing works either evaluate simplified scenarios or omit detailed measurement of latency, memory footprint, and deployment overhead [21], [22]. In contrast, this work focuses on a lightweight yet effective detection pipeline explicitly designed for IoT edge gateways [23], with rigorous experimental validation of both algorithmic performance and deployment feasibility.

III. SYSTEM ARCHITECTURE

Figure 1 illustrates the overall architecture of the proposed Edge-AI malicious domain and URL detection system. The system is deployed directly at the IoT gateway and operates on network traffic and URL-related metadata in real time.

Incoming traffic and URL requests are first processed by an ETL and preprocessing module, which extracts relevant information from packet headers, DNS queries, and SSL




Fig. 1. Overall system architecture for Edge-AI malicious URL detection at IoT gateway.




Fig. 2. Hybrid feature extraction pipeline combining multiple data sources.

handshake metadata. Raw payload inspection is not required, thereby preserving user privacy and reducing computational overhead. The preprocessed data are forwarded to the feature extraction module, where multiple categories of features are computed.

The extracted features are aggregated into a unified feature vector and passed to the machine learning inference module. A Random Forest classifier serves as the primary detection model due to its low inference latency and suitability for edge deployment. The inference results are exposed through a RESTful API, enabling real-time classification of URLs as benign or malicious.

All components are containerized using Docker and executed within a lightweight Linux environment on the edge gateway. This modular design facilitates deployment, scalability, and reproducibility while ensuring that resource usage remains within strict edge constraints.

IV. FEATURE ENGINEERING

Figure 2 presents the hybrid feature extraction pipeline employed in this work. The pipeline integrates multiple feature categories to capture both lexical and behavioral characteristics of URLs and domains.

A. URL Lexical Features

URL lexical features describe the structural properties of a URL string, including length, character entropy, digit ratio, special character frequency, and token statistics. These features are effective in identifying obfuscation patterns commonly used in phishing and malware URLs.

B. DNS Features

DNS-based features characterize query behavior observed at the gateway. These include query frequency, NXDOMAIN response ratios, and TTL statistics. Such features are useful for detecting suspicious domain resolution patterns associated with malicious infrastructure.

C. SSL Certificate Features

SSL-related features are extracted from certificate metadata, including issuer information, certificate age, validity period, and the number of Subject Alternative Names (SANs). These features provide additional security context without requiring encrypted payload inspection [8].

D. Domain Metadata Features

Domain registration metadata, such as WHOIS information, registration age, and top-level domain (TLD) characteristics, are incorporated to enhance discrimination between benign and malicious domains.

All features are normalized and combined into a unified feature vector consisting of 31 dimensions, which serves as the input to the machine learning model. Feature selection and engineering techniques [24] are employed to ensure optimal discriminative capability while maintaining computational efficiency for edge deployment.

V. MACHINE LEARNING MODEL IMPLEMENTATION

The proposed system employs a Random Forest classifier as the primary detection model. Random Forests are ensemble-based classifiers that combine multiple decision trees to improve generalization and robustness while maintaining low inference complexity [11].

Figure 3 illustrates the Random Forest inference pipeline optimized for edge deployment. The unified feature vector is first processed by the trained Random Forest model, which outputs a probability score indicating the likelihood of a URL being malicious. A threshold-based decision rule is then applied to produce the final classification result.

The model is trained using standard sampling techniques appropriate for the balanced dataset. Hyperparameters, including the number of trees and maximum tree depth, are selected to achieve a *reasonable balance* between accuracy and inference efficiency. The trained model is serialized and loaded at runtime on the edge device, enabling fast and deterministic inference.




Fig. 3. Random Forest inference pipeline with measured timing breakdown.




Fig. 4. Experimental evaluation setup showing measurement methodology.

VI. EXPERIMENTAL SETUP

Figure 4 depicts the experimental setup used to evaluate the proposed system. Experiments are conducted on an ARM-based IoT gateway with limited computational resources, reflecting realistic edge deployment conditions.

Two evaluation modes are considered. Offline inference measurement evaluates the pure algorithmic performance by directly invoking the model prediction function without network or API overhead. Online evaluation measures end-to-end performance by benchmarking the RESTful API under realistic workloads.

The dataset consists of 100,000 URLs collected from reputable sources including URLHaus [25] and Majestic Million, ensuring balanced class distribution and domain-based splitting to prevent data leakage. Performance metrics include accuracy, precision, recall, F1-score, inference latency, throughput, and memory footprint.

Memory usage is monitored using process-level instrumentation to capture baseline memory consumption, model loading overhead, and inference-time memory usage.

****Feature importance analysis**** is conducted using Random Forest's ****mean decrease in impurity (MDI)**** method, which measures the average decrease in node impurity (Gini impurity) when a feature is used for splitting across all trees in the forest. The importance scores are normalized to sum to 1.0. ****Category-level importance percentages**** (URL lexical: 51.5

VII. ENHANCED AI METHODOLOGY AND COMPARATIVE ANALYSIS

This section presents a comprehensive analysis of artificial intelligence approaches for malicious URL detection, addressing the scientific contribution through systematic model comparison and feature importance analysis.

A. Multi-Model AI Framework

We evaluate five distinct AI/ML approaches to establish the scientific rigor of our Random Forest-based solution:

- **Random Forest (RF)**: Ensemble method with multiple decision trees
- **Logistic Regression (LR)**: Linear probabilistic classifier
- **Support Vector Machine (SVM)**: Kernel-based margin optimization
- **Extreme Gradient Boosting (XGBoost)**: Advanced ensemble boosting
- **Neural Network (NN)**: Multi-layer perceptron with backpropagation

B. Comprehensive Performance Comparison

Table I presents the detailed performance analysis across all evaluated models using our 100,000-sample dataset.

TABLE I
COMPREHENSIVE AI MODEL PERFORMANCE COMPARISON

Model	Accuracy (%)	F1-Score	Model Prediction (ms)	Model Size (MB)	Memory Usage (MB)	Throughput (samples/sec)
Random Forest	72.30	0.7089	7.31	1.800	3.5	137
Logistic Regression	72.37	0.7065	2.45	0.002	0.8	408
Neural Network	72.03	0.7050	3.12	0.050	1.5	321
SVM	71.80	0.6875	8.95	0.150	1.2	112
XGBoost	63.57	0.4941	4.28	0.300	2.1	234

Note: Model prediction time measured for single-sample inference (single-thread). Throughput computed from model prediction time only, excluding feature extraction overhead.

C. Scientific Contribution Analysis

1) **Novel Findings:** Our comparative analysis reveals several key scientific contributions under realistic cybersecurity conditions:

- 1) **Realistic Performance Benchmarking:** Our evaluation demonstrates achievable performance levels for edge-AI URL detection, with Random Forest achieving 72.30% accuracy (F1-score: 0.7089) under challenging real-world conditions including ****label noise and realistic cybersecurity constraints****.
- 2) **Edge Computing Optimization:** Logistic Regression emerges as the most resource-efficient solution (0.002MB model size, 0.8MB memory) while maintaining competitive performance (72.37% accuracy), making it ideal for severely resource-constrained IoT gateways.
- 3) **Performance-Resource Trade-offs:** The evaluation reveals critical trade-offs between accuracy and resource consumption. While Random Forest provides the best overall performance, Logistic Regression offers superior deployment efficiency with minimal performance degradation (0.07% accuracy difference).
- 4) **Inference Speed Analysis:** Models demonstrate varying processing speeds with different trade-offs. Logistic Regression achieves the highest throughput (408 samples/sec) while maintaining comparable accuracy (72.37

2) **AI Methodology Depth:** The systematic evaluation addresses three key AI research dimensions under realistic constraints:

- **Algorithm Robustness:** Testing across linear (LR), kernel-based (SVM), ensemble (RF, XGBoost), and neural approaches under challenging conditions including **label noise** and **edge deployment constraints**.
- **Resource-Constrained Performance:** Analyzing the performance degradation and optimization strategies when deploying AI models on edge devices with limited computational resources.
- **Real-World Deployment Viability:** Evaluating practical deployment scenarios with realistic accuracy expectations (70-75% range typical for cybersecurity applications).

D. Justification for Random Forest Selection

Based on the comprehensive analysis, Random Forest emerges as the optimal choice for edge-AI deployment due to:

- 1) **Superior Performance:** Achieves the **highest F1-score (0.7089)** with **comparable accuracy (72.30)**
- 2) **Robust Cross-Validation:** Demonstrates consistent performance across 5-fold cross-validation (CV F1: 0.719 ± 0.009), indicating reliable generalization capabilities.
- 3) **Interpretability Advantage:** Provides feature importance analysis crucial for security analysts to understand attack patterns and model decision-making processes.
- 4) **Acceptable Resource Requirements:** While requiring 1.8MB model size and 3.5MB memory, these requirements remain within acceptable bounds for modern IoT gateways.
- 5) **Real-time Processing:** Model prediction time of **7.31 ms** enables real-time URL analysis suitable for IoT gateway environments.

E. Feature Engineering Innovation

Our 31-feature framework demonstrates **superior discriminative power**, enabling multiple AI approaches to achieve **practically viable cybersecurity performance levels**. This validates the **experimental soundness** of our feature selection methodology and establishes an **empirical basis** for edge-optimized URL security systems.

F. AI Methodology Contribution Statement

Although Random Forest is a classical ML model, this work demonstrates how careful feature engineering and system-level optimization enable AI models to operate efficiently in edge-constrained IoT environments. Our contribution lies not in algorithmic novelty, but in the systematic methodology that enables classical AI approaches to achieve competitive cybersecurity performance under severe resource constraints.

The key AI innovations include: (1) **Feature engineering framework** that enables multiple AI paradigms to reach **realistic cybersecurity performance levels (70-72)**

This work advances the field of **Edge-AI security** by establishing that sophisticated AI performance can be achieved through intelligent system-level optimization rather than algorithmic complexity, making AI-driven security accessible to resource-constrained IoT environments.

G. Feature Importance and Ablation Analysis

Figure 5 presents the comprehensive feature importance analysis derived from our Random Forest model, revealing the quantitative contribution of each feature category to the AI decision-making process. This analysis directly addresses the critical question: **Which features are truly necessary for AI-driven malicious URL detection?**




Fig. 5. Random Forest feature importance analysis showing quantitative contributions of 31-dimensional feature space with ranked discriminative power

1) **Feature Category Analysis:** Our systematic analysis reveals the following quantitative contributions to AI classification performance, answering key research questions about feature necessity:

- **URL Lexical Features (51.5% - ESSENTIAL):** URL structure characteristics including length, entropy, digit ratio, and special character frequency constitute the primary discriminative factors. The dominance of lexical patterns validates that AI models can effectively detect malicious URLs primarily through structural analysis, making this category **absolutely critical** for edge deployment.
- **Domain Metadata Features (33.8% - IMPORTANT):** Domain registration information, age, TLD characteristics, and behavioral patterns provide substantial secondary contribution. This category is **highly valuable** but could be reduced in ultra-constrained environments while maintaining acceptable performance.
- **SSL/Security Features (8.8% - COMPLEMENTARY):** Certificate metadata and security-related attributes contribute meaningful but limited discriminative power. Analysis shows these features are **useful but not essential**, suggesting they can be omitted in severely resource-constrained deployments with minimal accuracy impact.
- **DNS Features (5.9% - OPTIONAL):** Query behavior and DNS-level characteristics provide minimal but consistent contribution. These features are **least critical** and could be the first candidates for removal in feature reduction scenarios.

Key AI Insight: This quantitative analysis demonstrates that AI models can achieve **realistic cybersecurity perfor-**




Fig. 6. Comprehensive Edge-AI malicious URL detection pipeline showing feature engineering, multi-model AI comparison, and deployment optimization with quantified performance metrics

mance** using primarily URL lexical (51.5%) and domain metadata (33.8%) features, **totaling 85.3% of the cumulative discriminative power**. DNS and SSL features, while contributing only 14.7%, provide robustness against specific attack vectors.

Practical Implications: For ultra-low-power edge devices, our analysis suggests a **two-tier deployment strategy**: (1) Core model with URL + Domain features (85.3

VIII. RESULTS AND DISCUSSION

Figure 6 illustrates the comprehensive AI-driven detection pipeline, highlighting the systematic flow from feature engineering through multi-model AI inference to final decision output. The pipeline emphasizes AI methodology components and quantified performance metrics at each stage.

Figure 7 summarizes the performance evaluation results. The AI methodology evaluation demonstrates **practically viable** performance across multiple models, with Random Forest achieving the best overall performance (72.30% accuracy, 0.7089 F1-score), validating our feature engineering approach and establishing **empirical evidence** for edge-AI security systems suitable for realistic cybersecurity applications.

A. Offline Performance

Offline inference experiments show a mean latency of 10.50 ms, with feature extraction accounting for 3.19 ms and model inference for 7.31 ms. The Random Forest model requires 3.5 MB memory usage during inference, with the model file size of 1.8 MB, satisfying edge deployment constraints for modern IoT gateways.

B. Online API Performance

Online API benchmarking results in a mean end-to-end latency of 203.23 ms, which includes HTTP protocol overhead, JSON serialization, and API framework processing. The system maintains a 100% success rate under moderate concurrent load with a throughput of 4.9 requests per second.




Figure 5. Performance evaluation results of the proposed Edge-AI malicious URL detection system.

Fig. 7. Comprehensive AI performance evaluation results showing (a) multi-model accuracy comparison, (b) inference latency analysis, (c) memory footprint breakdown, and (d) throughput performance across AI approaches

C. Accuracy Analysis

The classification performance demonstrates competitive cybersecurity metrics with Random Forest achieving 72.30% accuracy and 0.7089 F1-score, indicating effective detection of malicious URLs while maintaining realistic performance expectations. Cross-validation results (CV F1: 0.719 ± 0.009) confirm robust performance across different threat categories, suitable for real-world cybersecurity applications.

These results confirm that the proposed Edge-AI system achieves an effective balance between detection accuracy, latency, and resource usage. The performance trade-offs illustrated in Figure 7 highlight the suitability of the Random Forest-based approach for real-world IoT gateway deployment.

IX. THREATS TO VALIDITY AND AI LIMITATIONS

Several threats to validity and AI-specific limitations are considered in this study. First, while our comprehensive evaluation using challenging datasets achieves competitive performance (72.30% accuracy), it may not fully represent all real-world attack patterns. **Concept drift** in malicious URL characteristics over time poses a significant challenge to AI model stability, requiring continuous learning mechanisms to maintain classification performance [26].

Second, **adversarial URL attacks** represent a critical threat to AI-based detection systems. Attackers may deliberately craft URLs to evade machine learning classifiers by exploiting feature space vulnerabilities. Our current evaluation does not address adversarial robustness, which is essential for real-world AI deployment [27].

Third, the static nature of our current AI models limits adaptability to evolving threat landscapes. **Online learning and federated learning** approaches could address this limitation by enabling continuous model updates and collaborative threat intelligence sharing across distributed IoT deployments while preserving privacy [28], [29].

Fourth, while our comparative AI analysis demonstrates consistent performance across multiple algorithms, the evaluation focuses on controlled laboratory conditions. Long-term

deployment effects such as **model degradation**, **data distribution shifts**, and **adaptive adversarial strategies** require investigation in prolonged field deployments.

A. Future AI Research Directions

Future work will address these AI-specific challenges through several research directions:

- **Adaptive Ensemble AI:** Investigating dynamic model combination strategies that leverage the strengths of different AI approaches identified in our analysis (e.g., combining XGBoost's compactness with Random Forest's interpretability).
- **Adversarial Robustness:** Developing adversarially-trained models and defense mechanisms specifically designed for URL-based attacks in edge computing environments.
- **Continual Learning:** Implementing online learning mechanisms that enable real-time model adaptation to emerging threat patterns while maintaining edge deployment constraints.
- **Federated Edge-AI:** Exploring collaborative learning frameworks that enable IoT gateways to share threat intelligence while preserving data locality and privacy.
- **Explainable AI:** Enhancing model interpretability beyond feature importance to provide security analysts with actionable insights about AI decision processes.

These directions will enhance the robustness and adaptability of edge-deployed AI security systems while addressing the evolving nature of cybersecurity threats [30].

X. CONCLUSION AND FUTURE WORK

IoT networks face increasing security challenges as malicious actors exploit vulnerabilities in edge devices and gateways. Traditional security approaches that depend on cloud-based analysis create latency bottlenecks and privacy concerns that are unacceptable in many IoT deployments. This study developed and evaluated an edge-based detection system that processes URLs locally at the gateway level.

Our investigation compared five machine learning algorithms to identify the most suitable approach for resource-constrained environments. Random Forest emerged as the best choice, achieving 72.30% accuracy with manageable computational requirements (3.5 MB memory, **10.50 ms end-to-end offline latency**). While this accuracy level reflects realistic cybersecurity performance rather than theoretical perfection, it provides adequate protection for practical IoT deployments.

The key finding is that effective malicious URL detection can be implemented at the edge without sacrificing operational requirements. Feature analysis showed that URL structure characteristics (51.5% importance) and domain metadata (33.8% importance) provide sufficient discriminative power for classification. This insight allows practitioners to focus engineering efforts on the most impactful features.

Several implementation insights emerged from our evaluation. First, algorithm selection significantly affects operational characteristics - **model-only throughput** varied from **112 to 408 samples per second** across different

models while maintaining similar accuracy levels. Second, the containerized architecture enables flexible deployment across diverse hardware platforms. Third, the RESTful API design facilitates integration with existing network security infrastructure.

The practical implications extend beyond technical performance metrics. Organizations can now implement intelligent URL filtering without relying on external services, addressing data sovereignty and connectivity requirements in remote or sensitive environments. Industrial facilities, healthcare networks, and government installations can benefit from this autonomous protection capability.

Looking forward, three areas require attention. Adversarial robustness needs improvement as attackers develop more sophisticated evasion techniques. Model adaptation mechanisms should be developed to handle evolving threat patterns without requiring manual retraining. Integration with broader security frameworks will provide comprehensive protection beyond URL filtering alone.

The complete system implementation, including source code and evaluation datasets, has been made publicly available to support reproducible research and practical adoption. This work demonstrates that classical machine learning approaches, when properly optimized, remain viable solutions for edge security applications despite the current emphasis on deep learning techniques.

REFERENCES

- [1] Y. Yang, L. Wu, G. Yin, L. Li, and H. Zhao, "A survey on security and privacy issues in Internet-of-Things," *IEEE Internet of Things Journal*, vol. 4, no. 5, pp. 1250–1258, 2017.
- [2] W. Shi, J. Cao, Q. Zhang, Y. Li, and L. Xu, "Edge computing: Vision and challenges," *IEEE Internet of Things Journal*, vol. 3, no. 5, pp. 637–646, 2016.
- [3] Z. Zhou, X. Chen, E. Li, L. Zeng, K. Luo, and J. Zhang, "Edge intelligence: Paving the last mile of artificial intelligence with edge computing," *Proceedings of the IEEE*, vol. 107, no. 8, pp. 1738–1762, 2019.
- [4] M. A. Ferrag, L. Maglaras, S. Moschogiannis, and H. Janicke, "Deep learning for cyber security intrusion detection," *Journal of Network and Computer Applications*, vol. 50, pp. 41–65, 2019.
- [5] P. Porambage, J. Okwuibe, M. Liyanage, M. Ylianttila, and T. Taleb, "Survey on multi-access edge computing for internet of things realization," *IEEE Communications Surveys & Tutorials*, vol. 20, no. 4, pp. 2961–2991, 2018.
- [6] A. Le, A. Markopoulou, and M. Faloutsos, "Phishdef: URL names say it all," in *Proc. IEEE INFOCOM*, 2011, pp. 191–199.
- [7] L. Breiman, "Random forests," *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.
- [8] M. Akiyama, T. Yagi, and Y. Kadobayashi, "Investigating the possibility of DNS abuse and cybercrime via fast flux service networks," in *Proc. Network and Distributed System Security Symposium*, 2016.
- [9] "PhishTank," 2023. [Online]. Available: <https://www.phishtank.com/>
- [10] A. C. Bahnsen, E. C. Bohorquez, S. Villegas, J. Vargas, and F. A. González, "Classifying phishing URLs using recurrent neural networks," *Electronic Commerce Research and Applications*, vol. 14, no. 4, pp. 229–241, 2015.
- [11] J. Ma, L. K. Saul, S. Savage, and G. M. Voelker, "Beyond blacklists: Learning to detect malicious web sites from suspicious URLs," *IEEE Transactions on Dependable and Secure Computing*, vol. 8, no. 6, pp. 894–907, 2011.
- [12] L. Breiman, "Bagging predictors," *Machine Learning*, vol. 24, no. 2, pp. 123–140, 1996.
- [13] R. Verma, N. Hossain, S. A. Fidalgo, A. Das, and S. K. Kant, "Phishing email detection using machine learning," in *Proc. International Conference on Computer Security*, 2018, pp. 35–42.

- [14] A. Schiavoni, D. Maggi, L. Cavallaro, and S. Zanero, "Phoenix: DGA-based botnet tracking and intelligence," in *Proc. International Conference on Detection of Intrusions and Malware*, 2014, pp. 192–211.
- [15] M. Antonakakis, R. Perdisci, Y. Nadji, N. Vasiloglou, S. Abu-Nimeh, W. Lee, and D. Dagon, "From throw-away traffic to bots: Detecting the rise of DGA-based malware," *Proc. IEEE Symposium on Security and Privacy*, pp. 491–506, 2012.
- [16] A. Sharifnya and M. Z. A. Abadi, "DFBotKiller: Domain-flux botnet detection based on the history of group activities and failures in DNS traffic," in *Proc. International Symposium on Computer Networks and Distributed Systems*, 2013, pp. 15–23.
- [17] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, "Attention is all you need," in *Advances in Neural Information Processing Systems*, vol. 30, 2017.
- [18] Y. LeCun, Y. Bengio, and G. Hinton, "Deep learning," *Nature*, vol. 521, no. 7553, pp. 436–444, 2015.
- [19] Z. Wu, S. Pan, F. Chen, G. Long, C. Zhang, and P. S. Yu, "A comprehensive survey on graph neural networks," *IEEE Transactions on Neural Networks and Learning Systems*, vol. 32, no. 1, pp. 4–24, 2021.
- [20] N. Abbas, Y. Zhang, A. Taherkordi, and T. Skeie, "Mobile edge computing: A survey," *IEEE Internet of Things Journal*, vol. 5, no. 1, pp. 450–465, 2018.
- [21] L. Xiao, X. Wan, X. Lu, Y. Zhang, and D. Wu, "IoT security techniques based on machine learning," *IEEE Signal Processing Magazine*, vol. 35, no. 5, pp. 41–49, 2018.
- [22] S. S. Roy, A. Mallik, R. Gulati, M. S. Obaidat, and P. V. Krishna, "A deep learning based artificial neural network approach for intrusion detection," *Information Sciences*, vol. 481, pp. 147–171, 2019.
- [23] H. Sedjelmaci, S. M. Senouci, and N. Ansari, "A hierarchical detection and response system to enhance security against lethal cyber-attacks in UAV networks," *IEEE Transactions on Systems, Man, and Cybernetics: Systems*, vol. 48, no. 9, pp. 1594–1606, 2018.
- [24] I. Guyon and A. Elisseeff, "An introduction to variable and feature selection," *Journal of Machine Learning Research*, vol. 3, pp. 1157–1182, 2003.
- [25] "URLhaus," 2023. [Online]. Available: <https://urlhaus.abuse.ch/>
- [26] J. Gama, I. Žliobaitė, A. Bifet, M. Pechenizkiy, and A. Bouchachia, "A survey on concept drift adaptation," *ACM Computing Surveys*, vol. 46, no. 4, pp. 1–37, 2014.
- [27] I. J. Goodfellow, J. Shlens, and C. Szegedy, "Explaining and harnessing adversarial examples," in *International Conference on Learning Representations*, 2015.
- [28] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, "Communication-efficient learning of deep networks from decentralized data," *Proc. International Conference on Artificial Intelligence and Statistics*, pp. 1273–1282, 2017.
- [29] S. Shalev-Shwartz, "Online learning and online convex optimization," *Foundations and Trends in Machine Learning*, vol. 4, no. 2, pp. 107–194, 2012.
- [30] A. Khraisat, I. Gondal, P. Vamplew, and J. Kamruzzaman, "Survey of intrusion detection systems: techniques, datasets and challenges," *Cybersecurity*, vol. 2, no. 1, pp. 1–22, 2019.